https://www.kaggle.com/datasets/uciml/red-wine-quality-cortez-et-al-2009

Importing Libraries

import numpy as np

import matplotlib.pyplot as plt

import pandas as pd
import seaborn as sns

from warnings import filterwarnings
filterwarnings(action='ignore')

Loading Dataset

wine = pd.read_csv("/content/winequality-red.csv")
wine.sample(25)

\overline{a}			
		-	7
	-	→	4

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	рН	sulphates	alcohol	quality
507	11.2	0.670	0.55	2.3	0.084	6.0	13.0	1.00000	3.17	0.71	9.5	6
547	10.6	0.310	0.49	2.5	0.067	6.0	21.0	0.99870	3.26	0.86	10.7	6
911	9.1	0.280	0.46	9.0	0.114	3.0	9.0	0.99901	3.18	0.60	10.9	6
1502	2 7.3	0.585	0.18	2.4	0.078	15.0	60.0	0.99638	3.31	0.54	9.8	5
6	7.9	0.600	0.06	1.6	0.069	15.0	59.0	0.99640	3.30	0.46	9.4	5
155	7.1	0.430	0.42	5.5	0.071	28.0	128.0	0.99730	3.42	0.71	10.5	5
1414	4 10.0	0.320	0.59	2.2	0.077	3.0	15.0	0.99940	3.20	0.78	9.6	5
833	11.6	0.470	0.44	1.6	0.147	36.0	51.0	0.99836	3.38	0.86	9.9	4
797	9.3	0.370	0.44	1.6	0.038	21.0	42.0	0.99526	3.24	0.81	10.8	7
1180	8.2	0.350	0.33	2.4	0.076	11.0	47.0	0.99599	3.27	0.81	11.0	6
1162	2 8.5	0.320	0.42	2.3	0.075	12.0	19.0	0.99434	3.14	0.71	11.8	7
555	15.5	0.645	0.49	4.2	0.095	10.0	23.0	1.00315	2.92	0.74	11.1	5
1562	2 7.2	0.695	0.13	2.0	0.076	12.0	20.0	0.99546	3.29	0.54	10.1	5
131	7.0	0.360	0.21	2.3	0.086	20.0	65.0	0.99558	3.40	0.54	10.1	6
562	9.0	0.540	0.49	2.9	0.094	41.0	110.0	0.99820	3.08	0.61	9.2	5
657	12.0	0.500	0.59	1.4	0.073	23.0	42.0	0.99800	2.92	0.68	10.5	7
564	13.0	0.470	0.49	4.3	0.085	6.0	47.0	1.00210	3.30	0.68	12.7	6
130	8.0	0.745	0.56	2.0	0.118	30.0	134.0	0.99680	3.24	0.66	9.4	5
1250	7.1	0.600	0.01	2.3	0.079	24.0	37.0	0.99514	3.40	0.61	10.9	6
110	5 6.3	0.570	0.28	2.1	0.048	13.0	49.0	0.99374	3.41	0.60	12.8	5
115	5 8.3	0.600	0.25	2.2	0.118	9.0	38.0	0.99616	3.15	0.53	9.8	5
1382	2 8.0	0.600	0.22	2.1	0.080	25.0	105.0	0.99613	3.30	0.49	9.9	5
1023	8.2	0.320	0.42	2.3	0.098	3.0	9.0	0.99506	3.27	0.55	12.3	6
346	6.6	0.815	0.02	2.7	0.072	17.0	34.0	0.99550	3.58	0.89	12.3	7

wine.info()

<class 'pandas.core.frame.DataFrame'>
 RangeIndex: 1599 entries, 0 to 1598
 Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64
4	chlorides	1599 non-null	float64
5	free sulfur dioxide	1599 non-null	float64

6	total sulfur dioxide	1599 non-null	float64
7	density	1599 non-null	float64
8	pH	1599 non-null	float64
9	sulphates	1599 non-null	float64
10	alcohol	1599 non-null	float64
11	quality	1599 non-null	int64

dtypes: float64(11), int64(1) memory usage: 150.0 KB

Description

wine.describe()



•	fixe acidit		citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	рН	sulphates
CC	ount 1599.00000	0 1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
m	ean 8.31963	7 0.527821	0.270976	2.538806	0.087467	15.874922	46.467792	0.996747	3.311113	0.658149
8	std 1.74109	6 0.179060	0.194801	1.409928	0.047065	10.460157	32.895324	0.001887	0.154386	0.169507
n	nin 4.60000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000	0.990070	2.740000	0.330000
2	5% 7.10000	0 0.390000	0.090000	1.900000	0.070000	7.000000	22.000000	0.995600	3.210000	0.550000
5	0% 7.90000	0 0.520000	0.260000	2.200000	0.079000	14.000000	38.000000	0.996750	3.310000	0.620000
7	5% 9.20000	0 0.640000	0.420000	2.600000	0.090000	21.000000	62.000000	0.997835	3.400000	0.730000
4										+

Finding Null Values

wine.isnull().sum()

	fixed acidity	0
	volatile acidity	0
	citric acid	0
	residual sugar	0
	chlorides	0
	free sulfur dioxide	0
	total sulfur dioxide	0
	density	0
	рН	0
	sulphates	0
	alcohol	0
	quality	0
	dtype: int64	

wine.groupby('quality').mean()

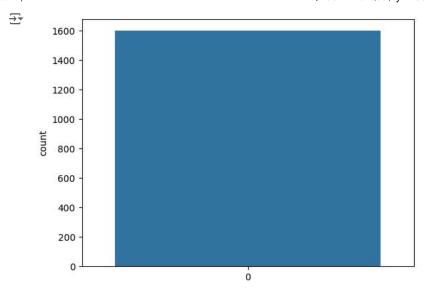
-	_	2	
	<u> </u>		

}		fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	рН	sulphates	alcohol
	quality											
	3	8.360000	0.884500	0.171000	2.635000	0.122500	11.000000	24.900000	0.997464	3.398000	0.570000	9.955000
	4	7.779245	0.693962	0.174151	2.694340	0.090679	12.264151	36.245283	0.996542	3.381509	0.596415	10.265094
	5	8.167254	0.577041	0.243686	2.528855	0.092736	16.983847	56.513950	0.997104	3.304949	0.620969	9.899706
	6	8.347179	0.497484	0.273824	2.477194	0.084956	15.711599	40.869906	0.996615	3.318072	0.675329	10.629519
	7	8.872362	0.403920	0.375176	2.720603	0.076588	14.045226	35.020101	0.996104	3.290754	0.741256	11.465913

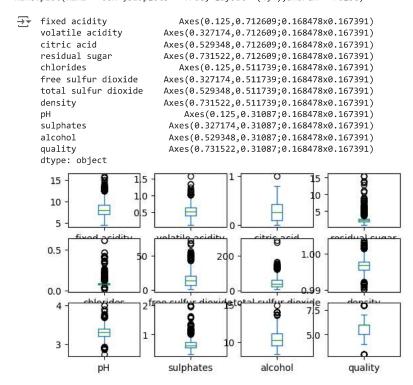
Data Analysis

Countplot:

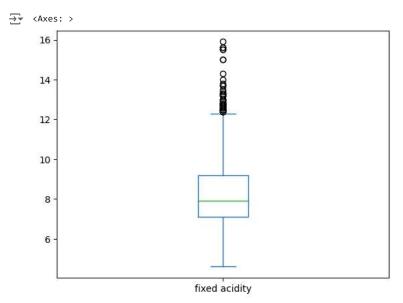
```
sns.countplot(wine['quality'])
plt.show()
```



wine.plot(kind ='box',subplots = True, layout =(4,4),sharex = False)

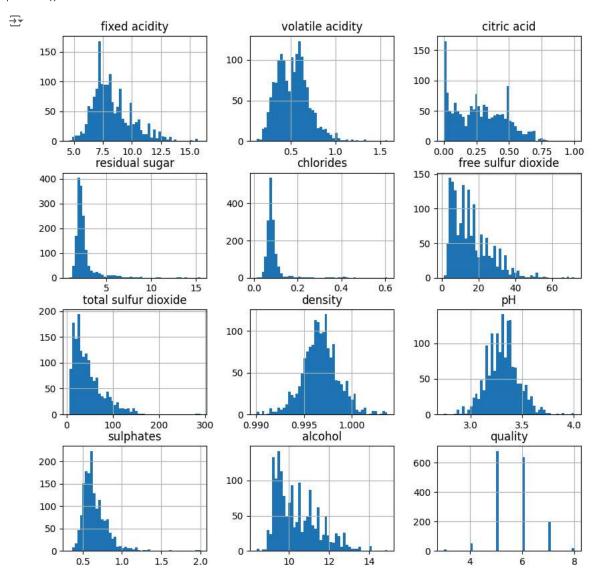


wine['fixed acidity'].plot(kind ='box')



Histogram

wine.hist(figsize=(10,10),bins=50)
plt.show()



Feature Selection

wine.sample(5)

		fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	рН	sulphates	alcohol	quality
	215	7.0	0.49	0.49	5.6	0.060	26.0	121.0	0.99740	3.34	0.76	10.5	5
	84	6.3	0.30	0.48	1.8	0.069	18.0	61.0	0.99590	3.44	0.78	10.3	6
	1208	7.2	0.36	0.46	2.1	0.074	24.0	44.0	0.99534	3.40	0.85	11.0	7
	1587	5.8	0.61	0.11	1.8	0.066	18.0	28.0	0.99483	3.55	0.66	10.9	6

wine['quality'].unique()

 \Rightarrow array([5, 6, 7, 4, 8, 3])

If wine quality is 7 or above then will consider as good quality wine wine['goodquality'] = [1 if $x \ge 7$ else 0 for x in wine['quality']] wine.sample(5)



•		fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	рН	sulphates	alcohol	quality	goodquality
	539	11.2	0.50	0.74	5.15	0.100	5.0	17.0	0.99960	3.22	0.62	11.2	5	0
	630	8.7	0.54	0.26	2.50	0.097	7.0	31.0	0.99760	3.27	0.60	9.3	6	0
	922	8.4	0.62	0.12	1.80	0.072	38.0	46.0	0.99504	3.38	0.89	11.8	6	0
	1365	7.8	0.50	0.09	2.20	0.115	10.0	42.0	0.99710	3.18	0.62	9.5	5	0

See total number of good vs bad wines samples
wine['goodquality'].value_counts()

Name: goodquality, dtype: int64

Separate depedent and indepedent variables

X = wine.drop(['quality','goodquality'], axis = 1)

Y = wine['goodquality']

Х

_		fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	рН	sulphates	alcohol
	0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
	1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8
	2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8
	3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8
	4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4
	1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5
	1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2
	1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0
	1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2
	1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0

1599 rows × 11 columns

print(Y)

\rightarrow	0	0
	1	0
	2	0
	3	0
	4	0
	1594	0
	1595	0
	1596	0
	1597	0
	1500	a

Name: goodquality, Length: 1599, dtype: int64

Feature Importance

```
from sklearn.ensemble import ExtraTreesClassifier
classifiern = ExtraTreesClassifier()
classifiern.fit(X,Y)
score = classifiern.feature_importances_
print(score)
```

[0.07642049 0.10357498 0.09302055 0.07172338 0.06930915 0.07053006 0.08113963 0.08497152 0.06915126 0.11090966 0.16924932]

Splitting Dataset

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.3,random_state=7)
```

Result

```
model_res=pd.DataFrame(columns=['Model', 'Score'])
```

LogisticRegression:

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train,Y_train)
y_pred = model.predict(X_test)

from sklearn.metrics import accuracy_score,confusion_matrix
# accuracy_score(Y_test,Y_pred)
model_res.loc[len(model_res)] = ['LogisticRegression', accuracy_score(Y_test,y_pred)]
model_res

Model Score

O LogisticRegression 0.86875
```

Using KNN:

```
from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier(n_neighbors=3)
model.fit(X_train,Y_train)
y_pred = model.predict(X_test)
from sklearn.metrics import accuracy_score
model_res.loc[len(model_res)] = ['KNeighborsClassifier', accuracy_score(Y_test,y_pred)]
model_res
<del>_</del>→
                     Mode1
                               Score
          LogisticRegression 0.870833
      1 KNeighborsClassifier 0.872917
  Using SVC:
from sklearn.svm import SVC
model = SVC()
model.fit(X_train,Y_train)
y_pred = model.predict(X_test)
from sklearn.metrics import accuracy_score
print("Accuracy Score:",accuracy_score(Y_test,y_pred))
model_res.loc[len(model_res)] = ['SVC', accuracy_score(Y_test,y_pred)]
model_res
Accuracy Score: 0.86875
                     Model
                               Score
         LogisticRegression 0.870833
      1 KNeighborsClassifier 0.872917
      2
                      SVC 0.868750
```

Using Decision Tree:

Using GaussianNB:

```
from sklearn.naive_bayes import GaussianNB
model3 = GaussianNB()
model3.fit(X_train,Y_train)
y_pred = model3.predict(X_test)
from sklearn.metrics import accuracy_score
print("Accuracy Score:",accuracy_score(Y_test,y_pred))
model_res.loc[len(model_res)] = ['GaussianNB', accuracy_score(Y_test,y_pred)]
model_res
Model
                             Score
          LogisticRegression 0.870833
         KNeighborsClassifier 0.872917
     2
                     SVC 0.868750
     3 DecisionTreeClassifier 0.864583
               GaussianNB 0.833333
```

Using Random Forest:

```
from sklearn.ensemble import RandomForestClassifier
model2 = RandomForestClassifier(random_state=1)
model2.fit(X_train, Y_train)
y_pred = model2.predict(X_test)
from sklearn.metrics import accuracy_score
print("Accuracy Score:",accuracy_score(Y_test,y_pred))
model_res.loc[len(model_res)] = ['RandomForestClassifier', accuracy_score(Y_test,y_pred)]
model\_res
Accuracy Score: 0.89375
                       Model
                                 Score
      0
             LogisticRegression 0.870833
      1
           KNeighborsClassifier 0.872917
      2
                         SVC 0.868750
          DecisionTreeClassifier 0.864583
                  GaussianNB 0.833333
      5 RandomForestClassifier 0.893750
```

!pip install xgboost

Using Xgboost:

```
import xgboost as xgb
model5 = xgb.XGBClassifier(random_state=1)
model5.fit(X_train, Y_train)
y_pred = model5.predict(X_test)
{\tt from \ sklearn.metrics \ import \ accuracy\_score}
print("Accuracy Score:",accuracy_score(Y_test,y_pred))
model_res.loc[len(model_res)] = ['XGBClassifier', accuracy_score(Y_test,y_pred)]
model_res
Accuracy Score: 0.891666666666667
                        Model
                                 Score
      0
             LogisticRegression 0.870833
      1
           KNeighborsClassifier 0.872917
      2
                         SVC 0.868750
      3
          DecisionTreeClassifier 0.864583
      4
                  GaussianNB 0.833333
      5 RandomForestClassifier 0.893750
                 XGBClassifier 0.891667
model_res = model_res.sort_values(by='Score', ascending=False)
model_res
₹
                       Model
                                 Score
      5 RandomForestClassifier 0.893750
      6
                 XGBClassifier 0.891667
      1
           KNeighborsClassifier 0.872917
             LogisticRegression 0.870833
      0
      2
                         SVC 0.868750
      3
          DecisionTreeClassifier 0.864583
      4
                  GaussianNB 0.833333
```

Start coding or generate with AI.